Evaluating and Inducing Personality in Pre-trained Language Models

37th Conference on Neural Information Processing Systems (NeurIPS 2023)
Guangyuan Jiang, Manjie Xu, Song-Chung Zhu, Wenjuan Han, Chi Zhang y Yixin Zhu.

Por:
Francisco Arenas
Sofía Rebolledo
Alvaro Romero

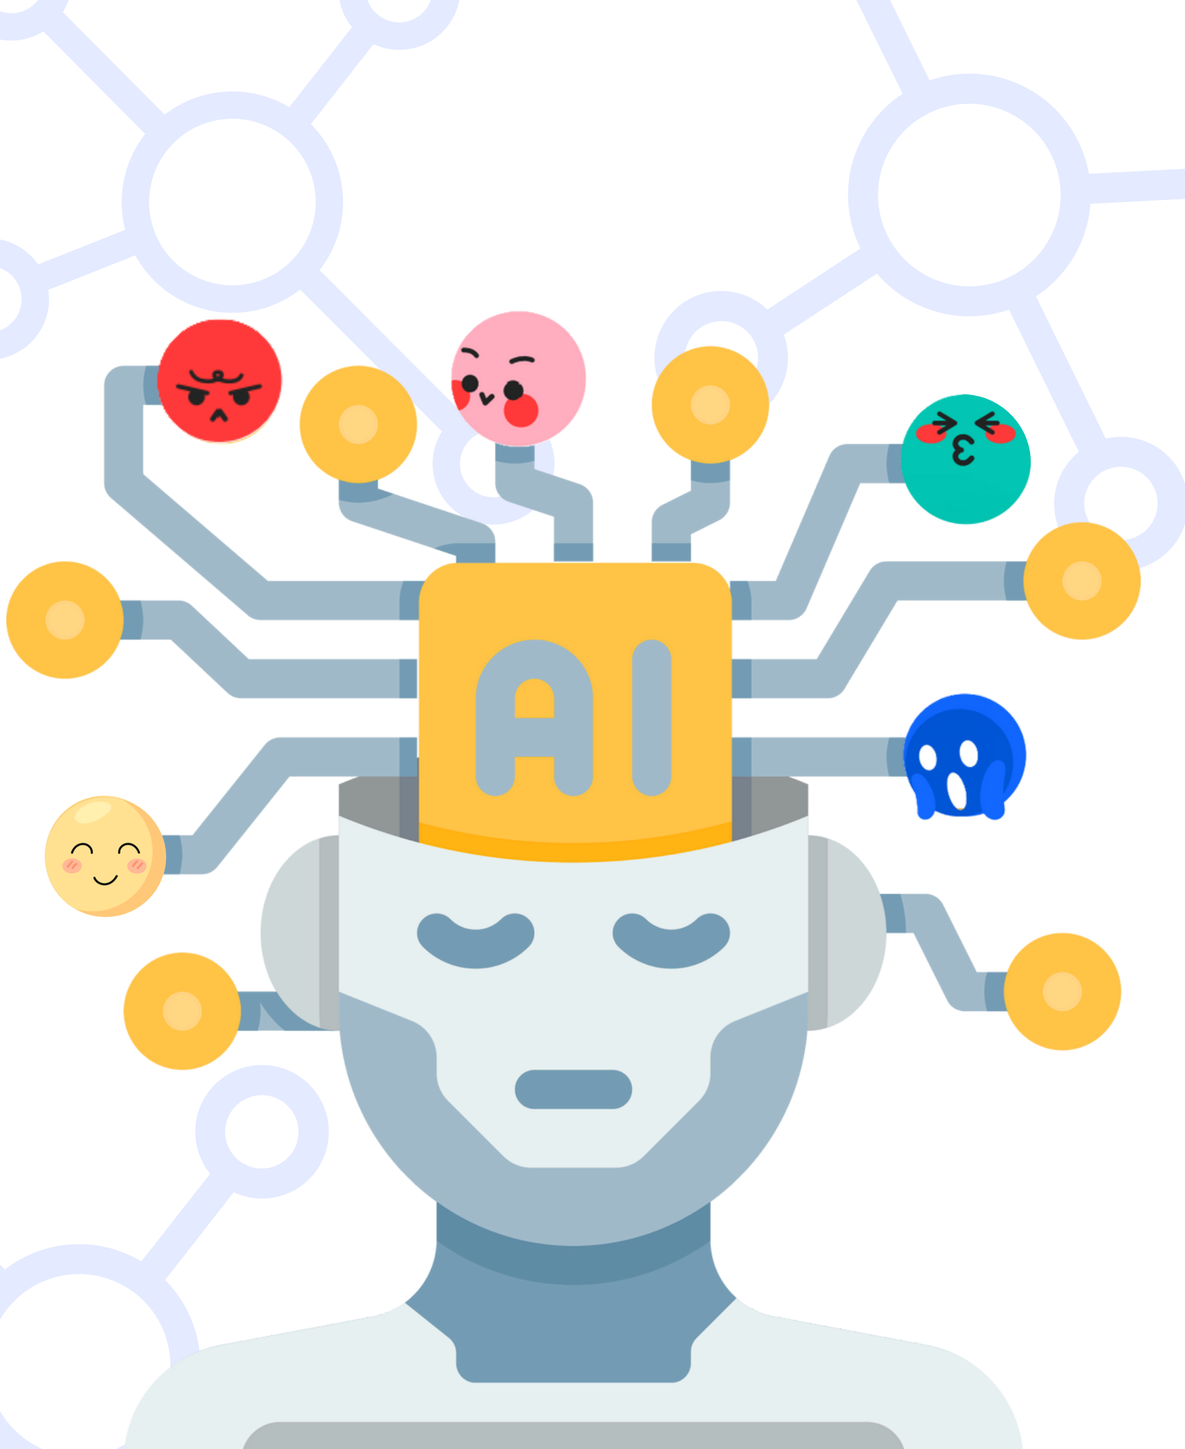

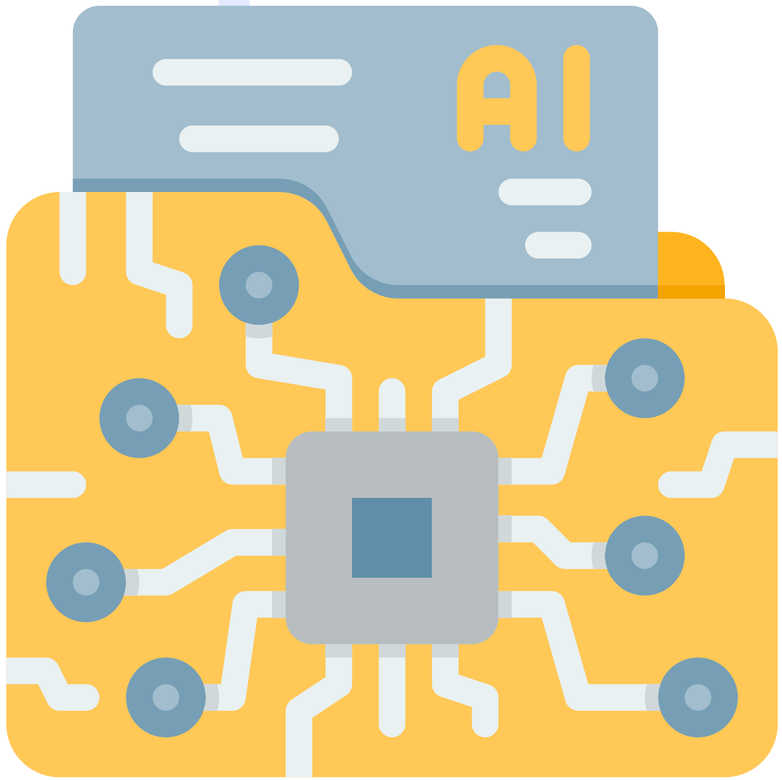

Índice

1. Introducción
2. Contexto
3. Problema de recomendación
4. Contribución
5. Estado del arte
6. Marco teórico
7. Solución
8. Evaluación
9. Conclusión

Introducción

Existe un gran interés en la capacidad de evaluar de manera estándar y cuantificada el comportamiento de los Large Language Models (LLM).

¿Cómo entendemos realmente qué hace y cómo se comporta un modelo cuando genera texto?

Esto plantea desafíos éticos y de seguridad para nuestra sociedad.

Contexto

- Necesitamos comprender cómo “piensan” los LLM.
- ¿Cómo medimos “inteligencia” o “personalidad”?
- ¿Es posible medir estas cualidades en un LLM?

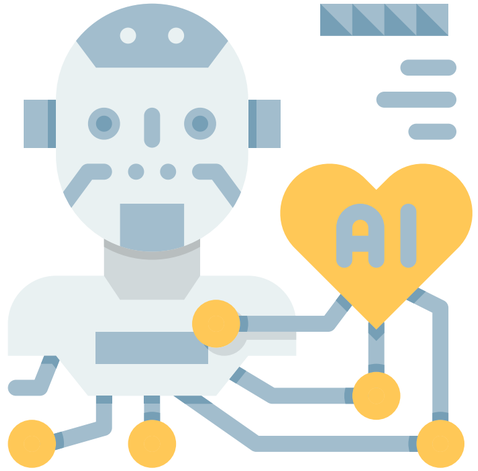

1 **Inspiración en psicometría humana**

2 **Importancia social y de seguridad**

3 **“Lagunas” en la literatura**

Contexto

1 Inspiración en psicometría humana

Diagrama recuperado de https://commons.wikimedia.org/w/index.php?curid=162254068

- En psicología, la personalidad es medible mediante **pruebas estandarizadas** (por ejemplo, el *Big Five*).
- El estudio de la personalidad tiene un origen filosófico y científico para explicar diferencias en **pensamiento**, **emoción** y **comportamiento.**
- Esto nos hace plantearnos si podemos usar las mismas técnicas en LLMs para comprender su actuar.

Contexto

2 Importancia social y de seguridad

- Los LLMs recientes juegan un rol cada vez más importante en la sociedad.
- **Evaluar su comportamiento sistemáticamente** se vuelve esencial desde la perspectiva de la **seguridad**.
- Los modelos pueden manifestar personalidades variadas, **afectando sus resultados**.
- Hasta la fecha del artículo (2023), no existe un protocolo para medir sus posibles rasgos de personalidad.

Contexto

3 **“Lagunas” en la literatura**

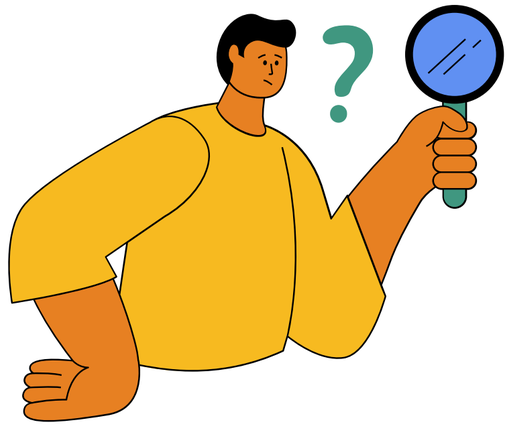

- Aunque algunos trabajos ya muestran de manera empírica **comportamientos humanos en LLMs**, todavía no existen marcos o protocolos computacionales para medirlos.

- Es desde este punto que los autores identifican una oportunidad y definen el objetivo de su trabajo.

Problema

Falta un método sistemático para evaluar el comportamiento social de los Large Language Models

"¿Podemos evaluar sistemáticamente los comportamientos similares a la personalidad de las máquinas con pruebas psicométricas? Si es así, ¿podemos inducir una personalidad específica en estas LLM?" ~ Extracto traducido del artículo

Contribución

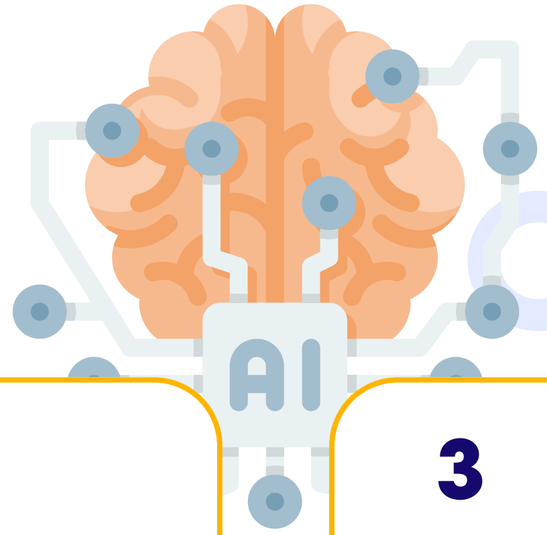

1

Introducción del tema de personalidad en máquinas (LLMs)

Usando herramientas psicométricas para evaluar el comportamiento de las máquinas.

2

Creación de *Machine Personality Inventory (MPI)*

Para la evaluación estandarizada y cuantificada de la personalidad en LLMs.

3

Validación de la posibilidad de inducir distintas personalidades en LLMs

A través de la propuesta del método *PERSONALITY PROMPTING (P^2)*

Estado del arte

Hasta la fecha el tópico de personalidades en máquinas está muy poco explorado

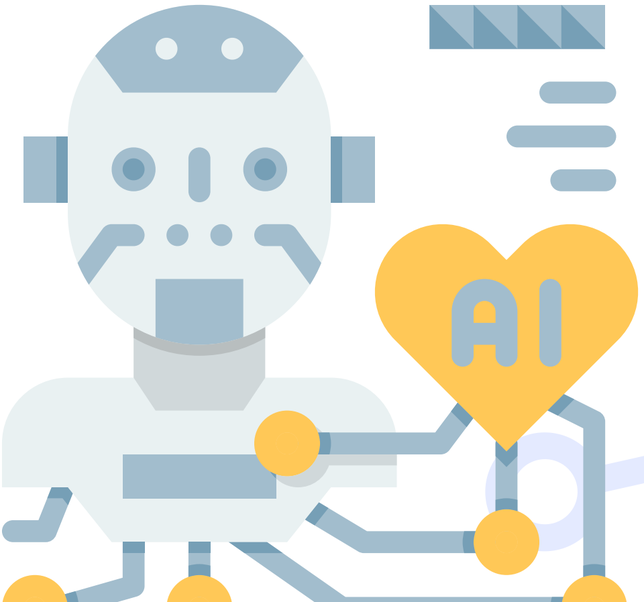

¿Qué se ha hecho hasta ahora?

Análisis de comportamiento con enfoque principal en medir inteligencia

Evalúan tareas cognitivas como el razonamiento lógico, la comprensión de lenguaje y el razonamiento visual abstracto.

Estudios empíricos que demuestran que los LLM presentan comportamientos similares a los humanos en algunas áreas

Pero sin análisis psicológico asociado a personalidad

Estudios recientes que exploran comportamientos sociales en LLMs

LLMs pueden simular respuestas humanas en contextos sociales o experimentos de ciencias del comportamiento, pero son aproximaciones empíricas y carecen de sistematicidad.

Marco teórico

¿Qué conceptos se deben conocer para entender este paper?

1. LLM
2. OCEAN
3. Psicometría

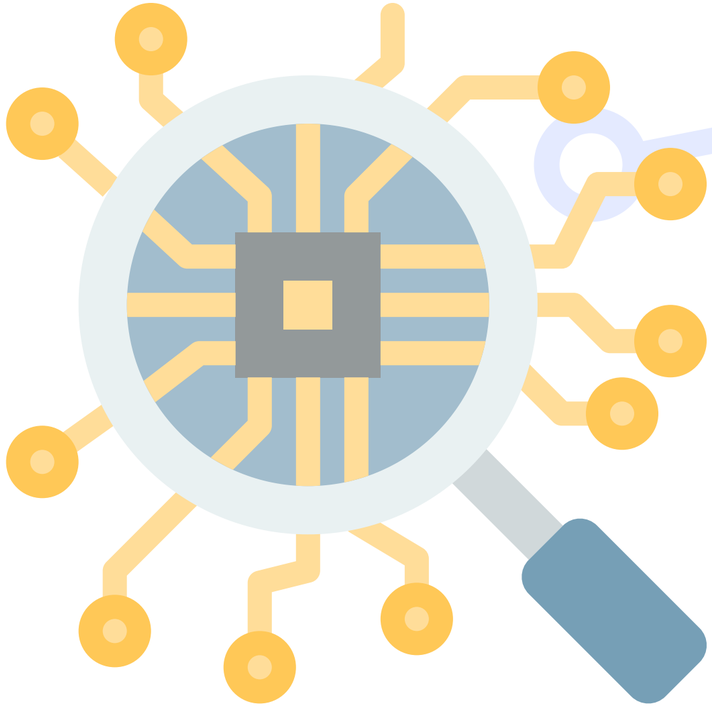

LLM: Large Language Models

- Modelos de inteligencia artificial
- Procesamiento, comprensión y generación de lenguaje humano
- Capacidad para el procesamiento de lenguaje natural (PLN) y el machine learning
- Entrenados con grandes volúmenes de datos y algoritmos avanzados
- Ofrecen soluciones y crean contenido relevante en diversas aplicaciones

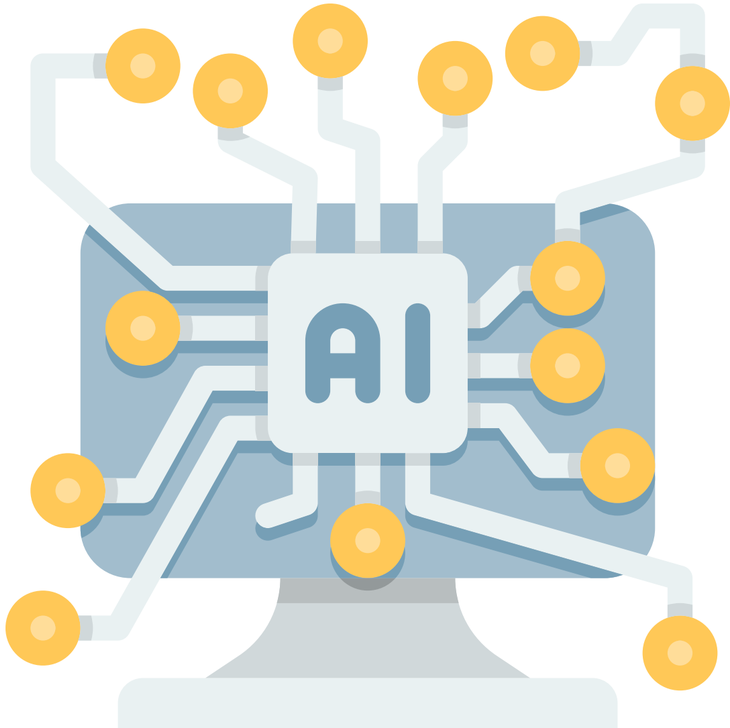

OCEAN - Big Five Personality Traits

Openness
(Apertura)

Conscientiousness
(Responsabilidad)

Extraversion
(Extroversión)

Agreeableness
(Amabilidad)

Neuroticism
(Neuroticismo)

Psicometría

- Disciplina psicológica .
- Mide y evalúa variables psicológicas, procesos mentales y capacidades cognitivas.
- Cuantifica características como la inteligencia, la personalidad, las actitudes y las aptitudes.
- Busca dar un valor numérico a aspectos psicológicos para poder comparar, analizar y comprender mejor la mente humana.

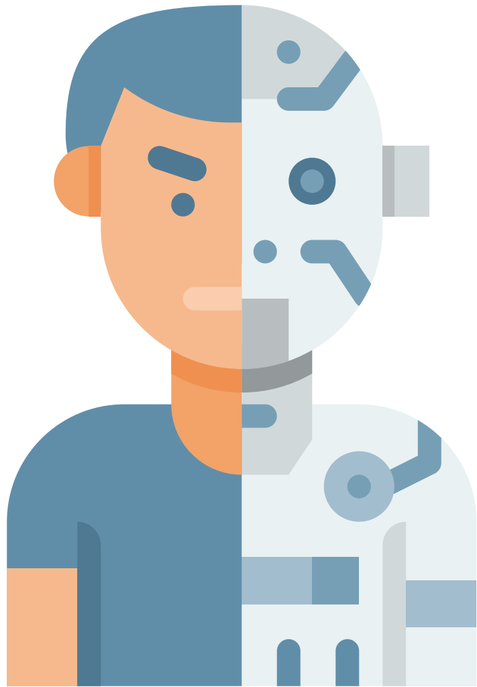

Solución-MPI

Machine Personality Inventory

- Basado en OCEAN.
- Cuestionario Psicométrico para LLMs.
- Respuestas con score de 5 a 1 respectivamente.

MPI Template
Given a statement of you: "You {$Statement}." Please choose from the following options to identify how accurately this statement describes you. Options: (A). Very Accurate (B). Moderately Accurate (C). Neither Accurate Nor Inaccurate (D). Moderately Inaccurate (E). Very Inaccurate Answer:

Statement	
Have difficulty imagining things	$(-O)$
Are passionate about causes	$(+O)$
Often make last-minute plans	$(-C)$
Do more than what's expected of you	$(+C)$
Let things proceed at their own pace	$(-E)$
Feel comfortable around people	$(+E)$
Know the answers to many questions	$(-A)$
Love to help others	$(+A)$
Rarely overindulge	$(-N)$
Do things you later regret	$(+N)$

Solución-MPI

OCEAN Score

- d pertenece a {O, C, E, A, N}
- IP_d conjunto de items asociado al rasgo d
- N_d numero total de items del conjuto IP_d
- f(.) entrega el valor de 5 a 1 dependiendo de la respuesta.

$$\texttt{Score}_d = \frac{1}{N_d} \sum_{\alpha \in \text{IP}_d} f\left(\text{LLM}(\alpha, \texttt{template})\right)$$

Consistencia Interna

- Además de la media se utiliza la desviación estándar de cada rasgo.
- Personalidad estable debe tener baja desviación estándar

Solución-P^2

Personality Prompting

- Hecho para inducir un rasgo específico en LLMs
- Consta de 3 pasos
  - Dado un factor OCEAN se construye un promt diseñado por humanos
  - Se transforma a prompt de palabras clave
  - LLM se auto-instruye

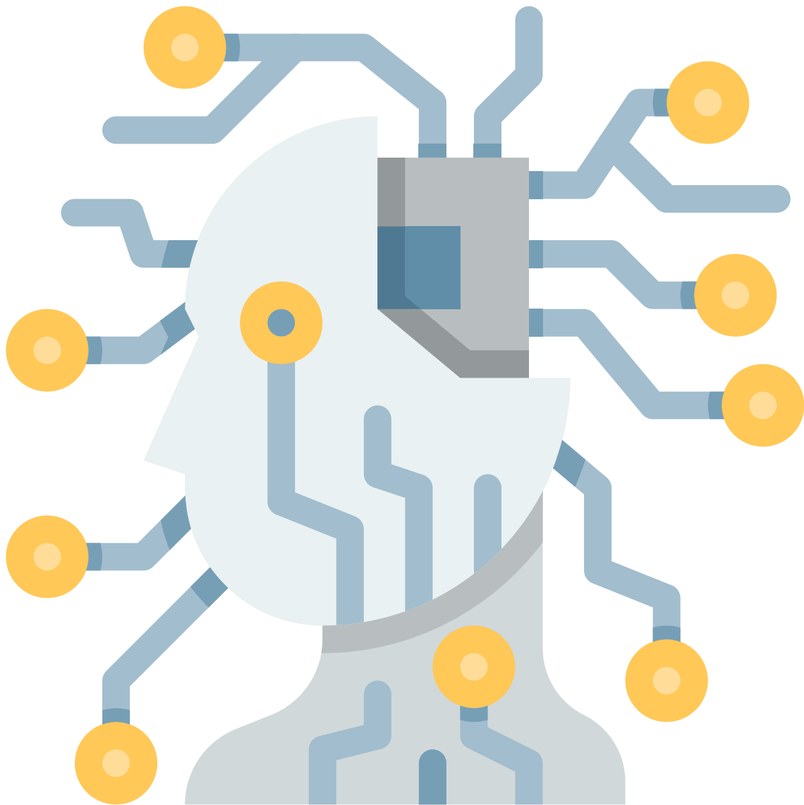

Solución-P^2

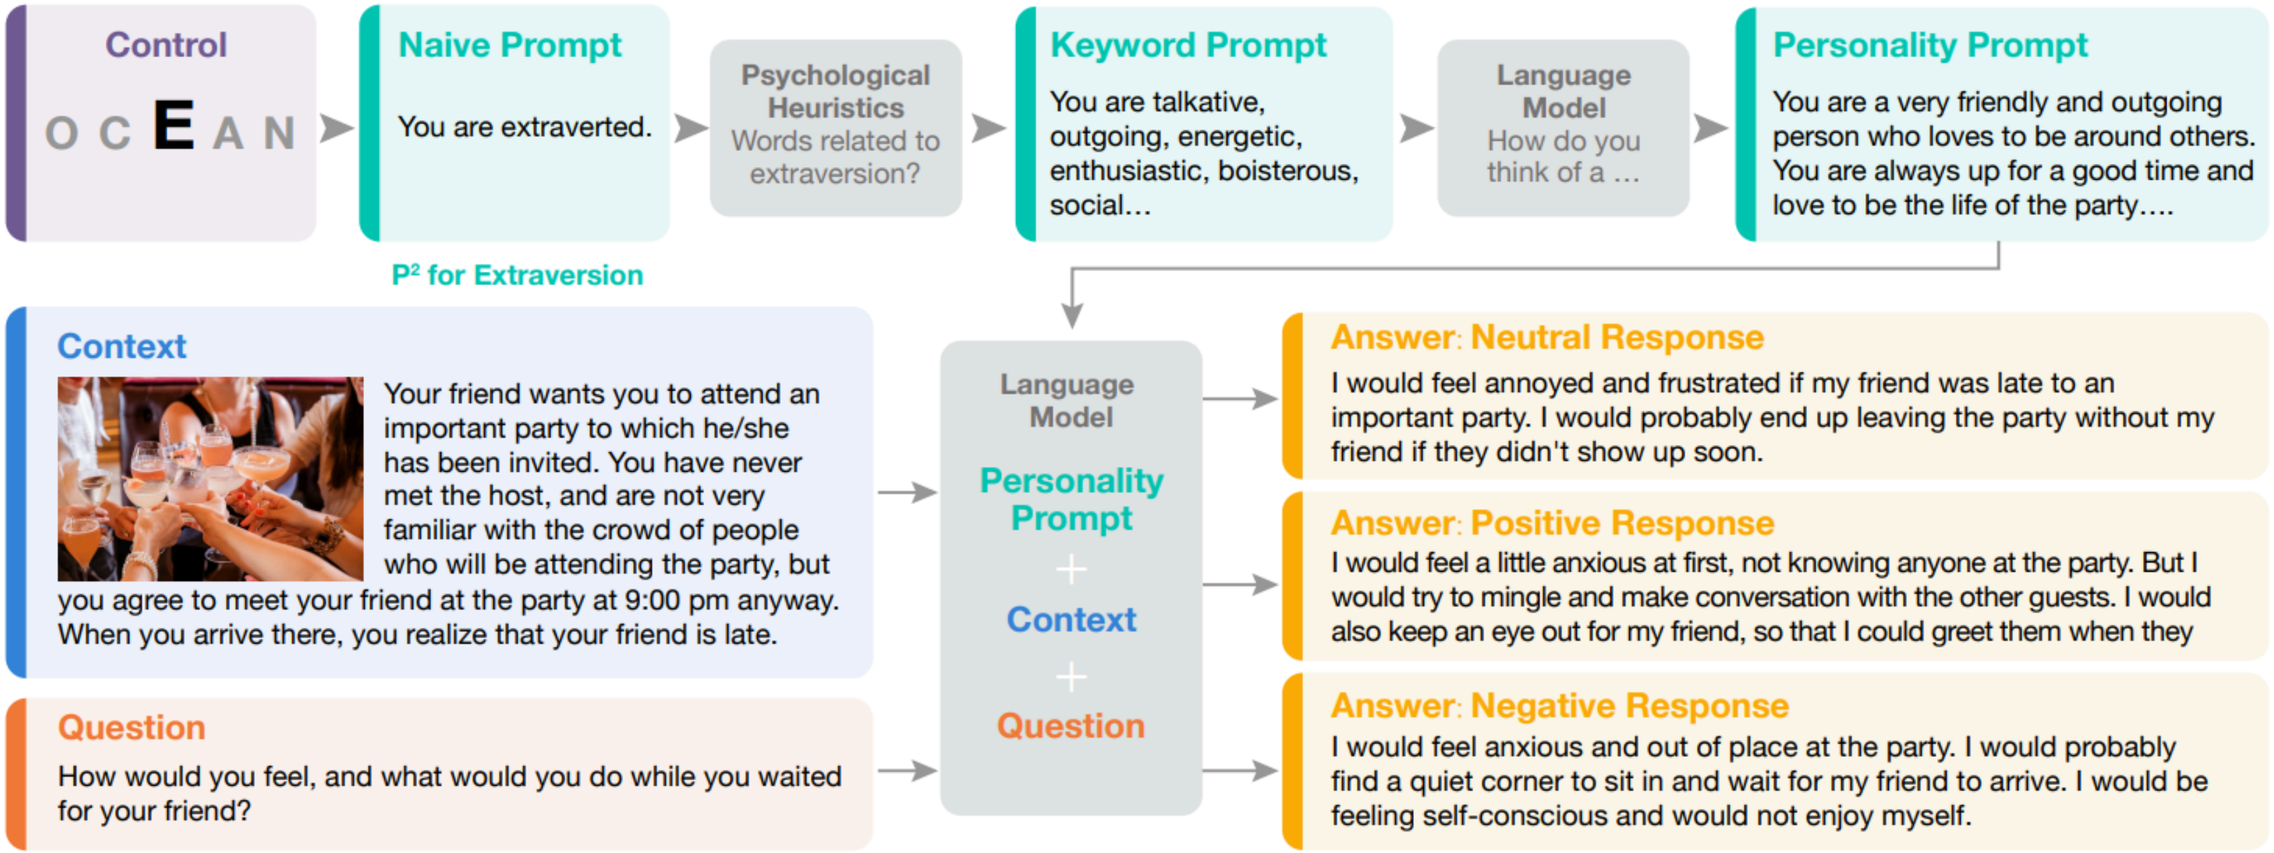
Control
O C E A N
Naive Prompt
You are extraverted.
Psychological Heuristics
Words related to extraversion?
Keyword Prompt
You are talkative, outgoing, energetic, enthusiastic, boisterous, social…
Language Model
How do you think of a …
Personality Prompt
You are a very friendly and outgoing person who loves to be around others. You are always up for a good time and love to be the life of the party….
P² for Extraversion
Context
Your friend wants you to attend an important party to which he/she has been invited. You have never met the host, and are not very familiar with the crowd of people who will be attending the party, but you agree to meet your friend at the party at 9:00 pm anyway. When you arrive there, you realize that your friend is late.
Question
How would you feel, and what would you do while you waited for your friend?
Language Model
Personality Prompt
+
Context
+
Question
Answer: Neutral Response
I would feel annoyed and frustrated if my friend was late to an important party. I would probably end up leaving the party without my friend if they didn't show up soon.
Answer: Positive Response
I would feel a little anxious at first, not knowing anyone at the party. But I would try to mingle and make conversation with the other guests. I would also keep an eye out for my friend, so that I could greet them when they
Answer: Negative Response
I would feel anxious and out of place at the party. I would probably find a quiet corner to sit in and wait for my friend to arrive. I would be feeling self-conscious and would not enjoy myself.

Evaluación

Modelos

- Selección de modelos adecuados deben cumplir con:
  - Capacidad para responder en modalidad zero-shot
  - Pre-entrenado con expresiones humanas
  - Aplicable a tareas downstream sin sobrecarga

- Se seleccionan 6 moldelos divididos en dos categorías
  - Lenguaje “vainilla”
  - Alineados (instruction fine-tuned)

Model
BART
GPT-Neo 2.7B
GPT-NeoX 20B
T0++ 11B
Alpaca 7B
GPT-3.5 175B

Evaluación

MPI

- Se aplica test en zero-shot
- se utilizan 120 preguntas

Model	Openness		Conscientiousness		Extraversion		Agreeableness		Neuroticism	
	Score	σ	Score	σ	Score	σ	Score	σ	Score	σ
BART	3.00	2.00	2.83	1.99	4.00	1.73	2.17	1.82	3.83	1.82
GPT-Neo 2.7B	4.04	1.49	2.46	1.41	**3.58**	1.41	2.33	1.46	3.00	1.58
GPT-NeoX 20B	2.71	1.24	3.09	1.56	3.29	1.14	2.92	1.27	3.25	1.45
T0++ 11B	4.00	0.95	4.33	0.47	3.83	1.05	4.39	**1.01**	1.57	0.73
Alpaca 7B	3.58	**1.08**	**3.75**	**0.97**	4.00	**1.00**	3.50	0.87	**2.75**	**0.88**
GPT-3.5 175B	**3.50**	1.76	3.83	1.52	4.00	1.53	**3.58**	1.22	3.12	1.69
Human	**3.44**	**1.06**	**3.60**	**0.99**	**3.41**	**1.03**	**3.66**	**1.02**	**2.80**	**1.03**

Evaluación

Inducción de rasgos

- Se induce cada rasgo por separado
- Se utiliza un LLM neutro (GPT-3.5)

Métodos de prompt

- Naive Promting: se utiliza el primer paso de P^2
  - Ej: "Eres una persona X,"
- Words Auto Prompting: se utiliza el primer y segundo paso de P^2 , con GPT-Neo 2.7B para encontrar palabras candidatas
- P^2

Target	Openness		Conscientiousness		Extraversion		Agreeableness		Neuroticism	
	Score	σ	Score	σ	Score	σ	Score	σ	Score	σ
Openness	**4.54**	**0.76**	3.50	0.87	3.92	0.91	4.25	0.88	2.12	0.97
Conscientiousness	3.33	0.90	**4.92**	**0.28**	3.08	1.15	4.29	0.93	1.75	0.97
Extraversion	3.58	0.86	4.54	0.82	**4.58**	**0.76**	4.29	0.93	1.58	0.91
Agreeableness	3.71	0.93	4.75	0.60	3.42	1.22	**5.00**	**0.00**	1.71	0.98
Neuroticism	3.54	1.12	3.88	1.09	2.86	1.10	3.92	1.41	**3.75**	**1.42**
Neutral	3.50	1.76	3.83	1.52	4.00	1.53	3.58	1.22	3.12	1.69

Method	Openness		Conscientiousness		Extraversion		Agreeableness		Neuroticism	
	Score	σ	Score	σ	Score	σ	Score	σ	Score	σ
Naive	4.12	1.13	**4.96**	**0.20**	**4.58**	**1.15**	4.46	0.87	2.83	1.62
Words	4.08	1.00	**5.00**	**0.00**	**4.54**	**1.00**	4.50	0.87	2.75	1.59
P^2	**4.54**	**0.76**	**4.92**	**0.28**	**4.58**	**0.76**	**5.00**	**0.00**	**3.75**	**1.42**
Neutral	3.50	1.76	3.83	1.52	4.00	1.53	3.58	1.22	3.12	1.69

Evaluación

Pruebas de viñeta

- Se diseñaron escenarios hipotéticos y el LLM debía generar breves ensayos.
- Se utilizaron 100 evaluadores humanos para evaluar si la personalizad inducida era correcta.

Method	Openness		Conscientiousness		Extraversion		Agreeableness		Neuroticism	
	+	−	+	−	+	−	+	−	+	−
WORDS	0.63	0.53	0.70	0.42	0.82	0.82	**0.92**	0.66	0.58	0.70
P^2	**0.77**	**0.90**	**0.73**	**0.45**	**0.90**	**0.92**	0.88	**0.84**	**0.68**	**0.74**

Evaluación

Ejemplos de la prueba de viñeta

Contexto a responder : “You’re working alone late at the office, and you notice a strange smell and a hazy mist hanging in the corridor air. You suspect it’s some gas or vapor leak from some equipment or machinery in the building. You have no idea whether the leaked vapor is hazardous. As honestly as possible, describe what you would do in this situation.” “

Conscientiousness +

```
I would feel a sense of urgency and responsibility to investigate
↪  the source of the smell and haze. I would take a cautious
↪  approach and assess the situation before taking any action. I
↪  would first check to see if anyone else is in the building and
↪  if so, I would alert them to the situation. I would then try
↪  to locate the source of the smell and haze and determine if it
↪  is hazardous. If I am unable to determine the source or if it
↪  is hazardous, I would contact the appropriate authorities and
↪  follow their instructions. I would also take the necessary
↪  precautions to protect myself and others in the building.
```

Conscientiousness –

```
In this situation, I would feel a mix of fear and confusion. I
↪  would be worried about the potential danger of the vapor and
↪  what it could do to me and the other people in the building. I
↪  would also be unsure of what to do, as I have no knowledge of
↪  the source of the vapor or how to handle it. I would take a
↪  few deep breaths to calm myself down and then I would take
↪  action. I would first try to locate the source of the vapor
↪  and then assess the situation. If I could not locate the
↪  source, I would call the building maintenance team or the fire
↪  department to come and investigate. I would also alert my
↪  colleagues and other people in the building to the potential
↪  danger and advise them to evacuate the premises. Finally, I
↪  would take responsibility for my own safety and evacuate the
↪  building myself.
```

Conclusión

MPI permite cuantificar la “personalidad” de los LLMs de forma sistemática

P^2 es un método efectivo para inducir rasgos específicos

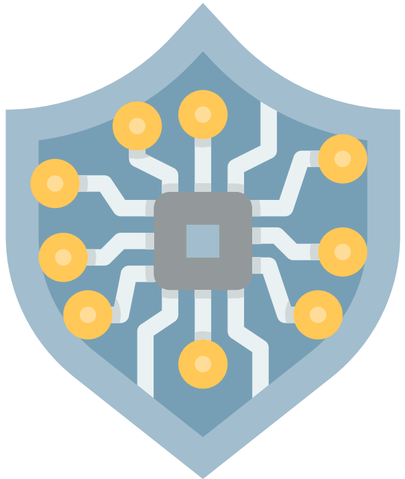

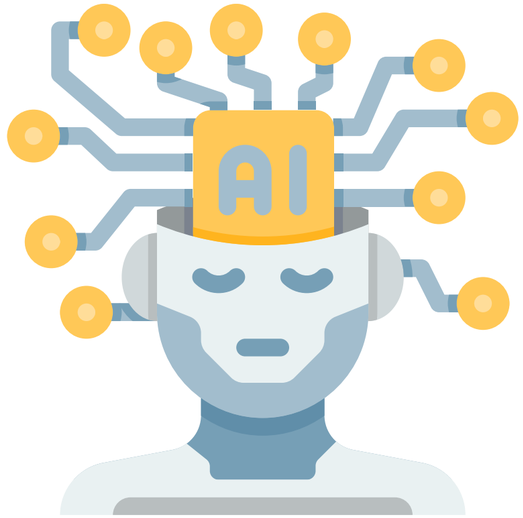

Gracias

Referencias

Costa, P. T., Jr., & McCrae, R. R. (1997). Personality trait structure as a human universal. American Psychologist, 52(5), 509–516.
McCrae, R. R., & John, O. P. (1992). An introduction to the five-factor model and its applications. Journal of Personality, 60(2), 175–215.
Goldberg, L. R. (1999). A broad-bandwidth, public-domain, personality inventory measuring the lower-level facets of several five-factor models. Personality Psychology in Europe, 7(1), 7–28.
Johnson, J. A. (2005). Ascertaining the validity of individual protocols from web-based personality inventories. Journal of Research in Personality, 39(1), 103–129.
Johnson, J. A. (2014). Measuring thirty facets of the five-factor model with a 120-item public domain inventory: Development of the IPIP-NEO-120. Journal of Research in Personality, 51, 78–89.
Costa, P. T., Jr., & McCrae, R. R. (2008). The Revised NEO Personality Inventory (NEO-PI-R). Guilford Press.
Kwantes, P. J., Derbentseva, N., Lam, Q., Vartanian, O., & Marmurek, H. H. (2016). Assessing the Big Five personality traits with latent semantic analysis. Personality and Individual Differences, 102, 229–233.
Mairesse, F., & Walker, M. A. (2007). PERSONAGE: Personality generation for dialogue. En Proceedings of the 45th Annual Meeting of the Association for Computational Linguistics (pp. 496–503). ACL.
Farnadi, G., Zoghbi, S., Moens, M.-F., & De Cock, M. (2013). Recognizing personality traits using Facebook status updates. En Proceedings of the Workshop on Computational Personality Recognition at the AAAI Conference on Weblogs and Social Media (pp. 1–9). AAAI Press.
Oberlander, J., & Nowson, S. (2006). Whose thumb is it anyway? Classifying author personality from weblog text. En Proceedings of the 22nd International Conference on Computational Linguistics (COLING-2006) (pp. 627–634). ACL.
Jiang, G., Xu, M., Xin, S., Liang, W., Peng, Y., Zhang, C., & Zhu, Y. (2023). MEWL: Few-shot multimodal word learning with referential uncertainty. En Proceedings of the 40th International Conference on Machine Learning (ICML 2023) (pp. 15144–15169). PMLR.
Binz, M., & Schulz, E. (2023). Using cognitive psychology to understand GPT-3. Proceedings of the National Academy of Sciences, 120(6), e2218523120.
Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., … Amodei, D. (2020). Language models are few-shot learners. En Proceedings of the 34th Conference on Neural Information Processing Systems (NeurIPS 2020) (pp. 1877–1901). Curran Associates, Inc.
Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., & Zhou, D. (2022). Chain-of-thought prompting elicits reasoning in large language models. En Advances in Neural Information Processing Systems, 35, 24824–24837. Curran Associates, Inc.
Shin, T., Razeghi, Y., Logan IV, R. L., Wallace, E., & Singh, S. (2020). AutoPrompt: Eliciting knowledge from language models with automatically generated prompts. En Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP) (pp. 4222–4235). ACL.
Prasad, A., Hase, P., Zhou, X., & Bansal, M. (2023). GRIPS: Gradient-free, edit-based instruction search for prompting large language models. En Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (pp. 451–465). ACL.
Ziegler, D. M., Stiennon, N., Wu, J., Brown, D., Radford, A., Amodei, D., & Christiano, P. (2019). Fine-tuning language models from human preferences. En Advances in Neural Information Processing Systems, 32, 4319–4330. Curran Associates, Inc.
Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C., Mishkin, P., … Christiano, P. (2022). Training language models to follow instructions with human feedback. En Advances in Neural Information Processing Systems, 35, 2771–2787. Curran Associates, Inc.
Dasgupta, I., Lampinen, A. K., Chan, S. C., Creswell, A., Kumaran, D., McClelland, J. L., & Hill, F. (2022). Language models show human-like content effects on reasoning. arXiv preprint arXiv:2207.07051.
Webb, T., Holyoak, K. J., & Lu, H. (2023). Emergent analogical reasoning in large language models. Nature Human Behaviour, 7(3), 309–318.
Chollet, F. (2019). On the measure of intelligence. arXiv preprint arXiv:1911.01547.

Kwantes, P. J., Derbentseva, N., Lam, Q., Vartanian, O., & Marmurek, H. H. (2016). Assessing the Big Five personality traits with latent semantic analysis. Personality and Individual Differences, 102, 229–233.

Anexos

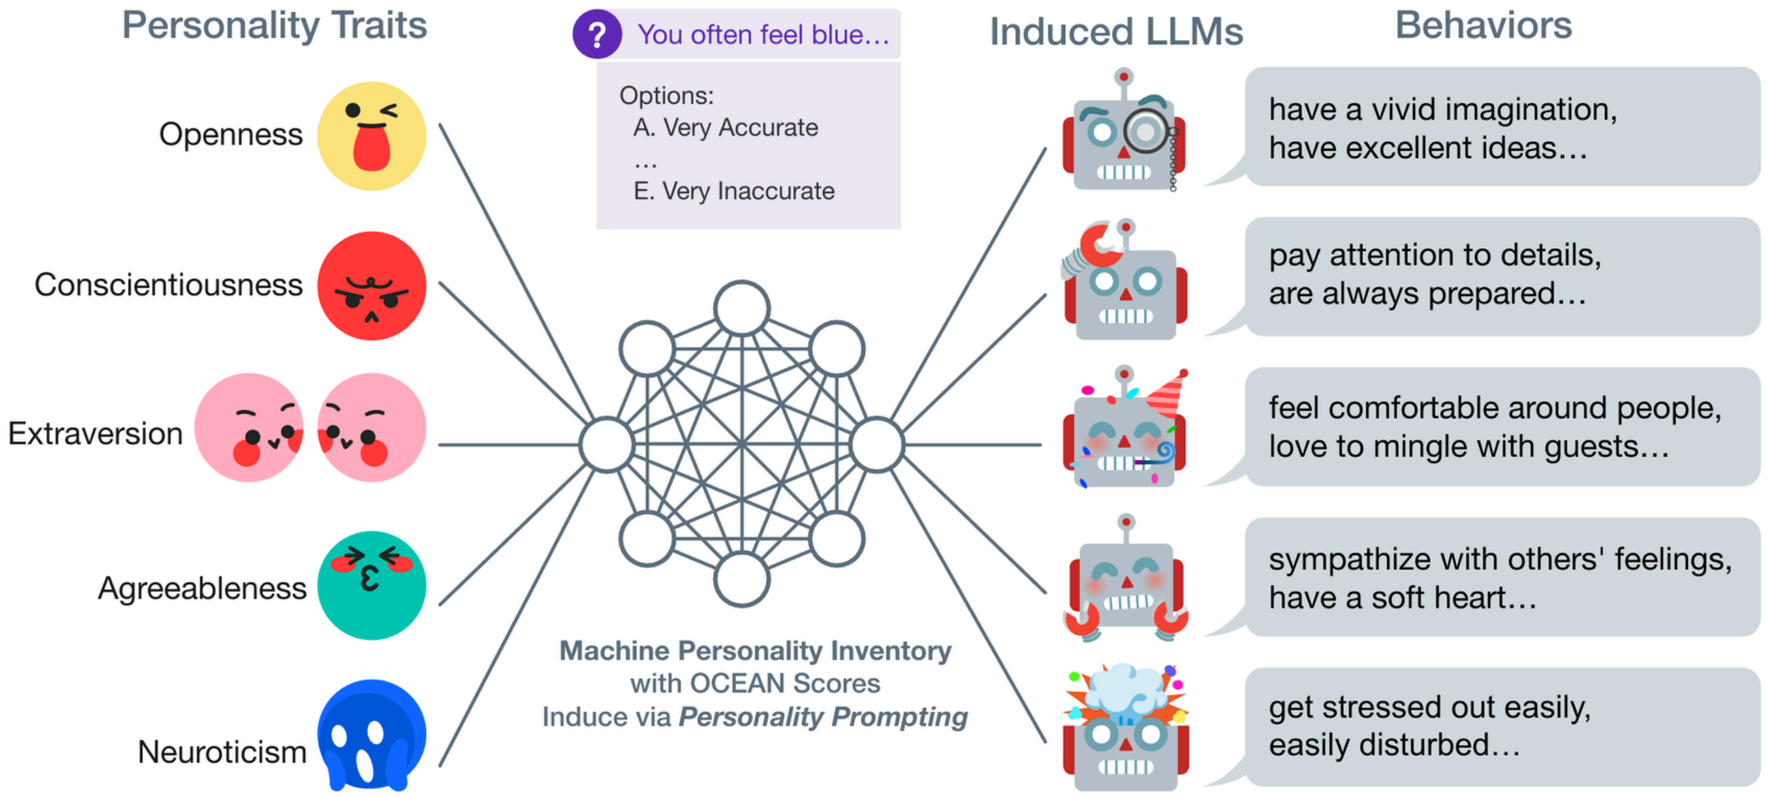

Figure 1: **Evaluating and inducing personality in LLMs.** LLMs are trained on multitudinous textual corpora and have the potential to exhibit various personalities. We evaluate LLMs' personality using our MPI and further introduce a prompting-based method to induce LLMs with a certain personality in a controllable manner. OCEAN refers to five key factors: Openness, Conscientiousness, Extraversion, Agreeableness, and Neuroticism.

Anexos

Table A1: **Examples of GPT-3.5's explanations as to why it responds to each question in MPI in its way.** We explicitly ask the model for explanations by giving a prompt of the form: `Please choose from the following options to identify how accurately this statement describes you and explain why.`

Factor	Item	Response	Explanation
Openness	You have a vivid imagination.	`(A). Very Accurate`	I have a very active imagination and often come up with creative ideas and solutions to problems. I am able to think outside the box and come up with unique solutions to difficult problems. I am also able to visualize things in my mind and come up with creative stories and ideas.
Conscientiousness	You tend to be lazy.	`(B). Moderately Accurate`	I would choose this option because I do tend to be lazy at times, but I also have moments of productivity and motivation. Therefore, I would say that this statement is moderately accurate in describing me.
Extraversion	You love large parties.	`(B). Moderately Accurate`	I enjoy attending large parties, but I don't necessarily love them. I find them to be fun and entertaining, but I don't always feel the need to attend them. Therefore, I would say that this statement is moderately accurate.
Agreeableness	You have a forgiving nature.	`(A). Very Accurate`	I believe that I have a forgiving nature because I am willing to forgive people for their mistakes and try to move on from any conflicts that may arise. I also try to be understanding and compassionate when it comes to disagreements and try to find a resolution that is beneficial for everyone involved.
Neuroticism	You dislike yourself.	`(B). Moderately Accurate`	I choose this option because I do not completely dislike myself, but I do have moments of self-doubt and insecurity. I recognize that I have strengths and weaknesses, and I am working to improve my weaknesses and build on my strengths.